



AMITY UNIVERSITY NOIDA
AMITY SCHOOL OF ENGINEERING AND TECHNOLOGY
ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING
PRACTICAL FILE

Made By-
Shourya Solanki
A2305223569
B.Tech 4CSE-3Y

Submitted To-
Dr Pintu Kumar Ram

PRACTICAL – 1

AIM:

- 1) To read the Iris Dataset.
- 2) To find any missing values.
- 3) To find the average value of each column.
- 4) To find the maximum value in each column.
- 5) To find the minimum value in each column.

LIBRARIES IMPORTED: Pandas

INPUT:

- 1) Import pandas as pd

```
>>>
```

```
>>> #Correct file path to the CSV file
```

```
>>> file_path = r'C:\Users\Peter\Downloads\Iris.csv'
```

```
>>>
```

```
>>> #Read the entire CSV file into a DataFrame
```

```
>>> iris_df = pd.read_csv (file_path)
```

```
>>>
```

```
>>> #Display the entire dataset (if its not too large)
```

```
>>> print(iris_df)
```

OUTPUT:

```
   Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species
0    1             5.1           3.5           1.4           0.2  Iris-setosa
1    2             4.9           3.0           1.4           0.2  Iris-setosa
2    3             4.7           3.2           1.3           0.2  Iris-setosa
3    4             4.6           3.1           1.5           0.2  Iris-setosa
4    5             5.0           3.6           1.4           0.2  Iris-setosa
..   ...           ...           ...           ...           ...   ...
145 146             6.7           3.0           5.2           2.3  Iris-virginica
146 147             6.3           2.5           5.0           1.9  Iris-virginica
147 148             6.5           3.0           5.2           2.0  Iris-virginica
148 149             6.2           3.4           5.4           2.3  Iris-virginica
149 150             5.9           3.0           5.1           1.8  Iris-virginica

[150 rows x 6 columns]
```

2) import pandas as pd

```
>>>
```

```
>>> #File path to the CSV file
```

```
>>> file_path = r'C:\Users\Peter\Downloads\Iris.csv'
```

```
>>>
```

```
>>> #Read the CSV file into a DataFrame
```

```
>>> iris_df = pd.read_csv(file_path)
```

```
>>>
```

```
>>> #Check for missing values in each column
```

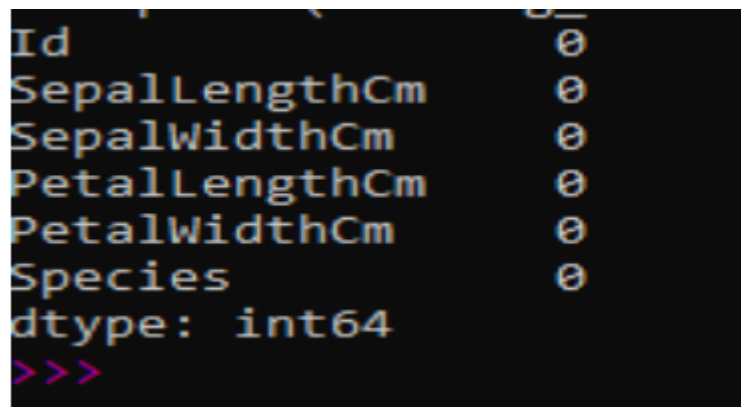
```
>>> missing_values = iris_df.isna().sum()
```

```
>>>
```

```
>>> #Display missing values count for each column
```

```
>>> print(missing_values)
```

OUTPUT:

A screenshot of a terminal window with a black background and yellow text. It displays the output of the print statement from the previous code block. The output shows a Series with six columns: Id, SepalLengthCm, SepalWidthCm, PetalLengthCm, PetalWidthCm, and Species. Each column has a value of 0, indicating no missing values. Below the Species column, it shows 'dtype: int64' and a prompt '>>>' in red.

```
Id      0
SepalLengthCm  0
SepalWidthCm   0
PetalLengthCm  0
PetalWidthCm   0
Species      0
dtype: int64
>>>
```

3) # Select only numeric columns from the dataset

```
>>> numeric_columns = iris_df.select_dtypes(include=['float64', 'int64'])
```

```
>>>
```

```
>>> # Calculate the mean (average) of each numeric column
```

```
>>> average_values = numeric_columns.mean()
```

```
>>>
```

```
>>> # Display the average for each numeric column
```

```
>>> print(average_values)
```

OUTPUT:

```
Id          75.500000
SepalLengthCm  5.843333
SepalWidthCm   3.054000
PetalLengthCm  3.758667
PetalWidthCm   1.198667
dtype: float64
>>>
```

4) # Select only numeric columns from the dataset

```
>>> numeric_columns = iris_df.select_dtypes(include=['float64', 'int64'])
```

```
>>>
```

```
>>> # Find the maximum value of each numeric column
```

```
>>> max_values = numeric_columns.max()
```

```
>>>
```

```
>>> # Display the maximum value for each numeric column
```

```
>>> print(max_values)
```

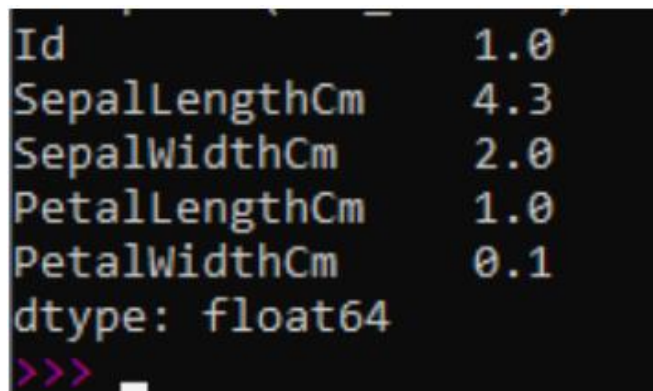
OUTPUT:

```
>>> print(max_values)
Id          150.0
SepalLengthCm   7.9
SepalWidthCm   4.4
PetalLengthCm   6.9
PetalWidthCm   2.5
dtype: float64
>>>
```

5) # Select only numeric columns from the dataset

```
>>> numeric_columns = iris_df.select_dtypes(include=['float64', 'int64'])
>>>
>>> # Find the minimum value of each numeric column
>>> min_values = numeric_columns.min()
>>>
>>> # Display the minimum value for each numeric column
>>> print(min_values)
```

OUTPUT:



```
Id      1.0
SepalLengthCm  4.3
SepalWidthCm   2.0
PetalLengthCm  1.0
PetalWidthCm   0.1
dtype: float64
>>>
```

PRACTICAL – 2

AIM:

Implementing k-nearest neighbour algorithm in python.

INPUT:

Import matplotlib.pyplot as plt

From sklearn.neighbors import KNeighborsClassifier as knc

#dataset

x = [4, 5, 10, 4, 3, 11, 14, 8, 10, 12]

y = [21, 19, 24, 17, 16, 25, 24, 22, 21, 21]

classes = [0, 0, 1, 0, 0, 1, 1, 0, 1, 1]

plt.plot(x, y, '*')

plt.show()

printf('Plotting the given DataSet')

plt.scatter(x = x, y = y, c = classes, marker = '*')

plt.show()

data = list(zip(x, y))

printf('Making the Array')

print(data)

k = knc(n_neighbors = 3)

k.fit(data, classes)

new_x = 8

new_y = 21

np = [(new_x, new_y)]

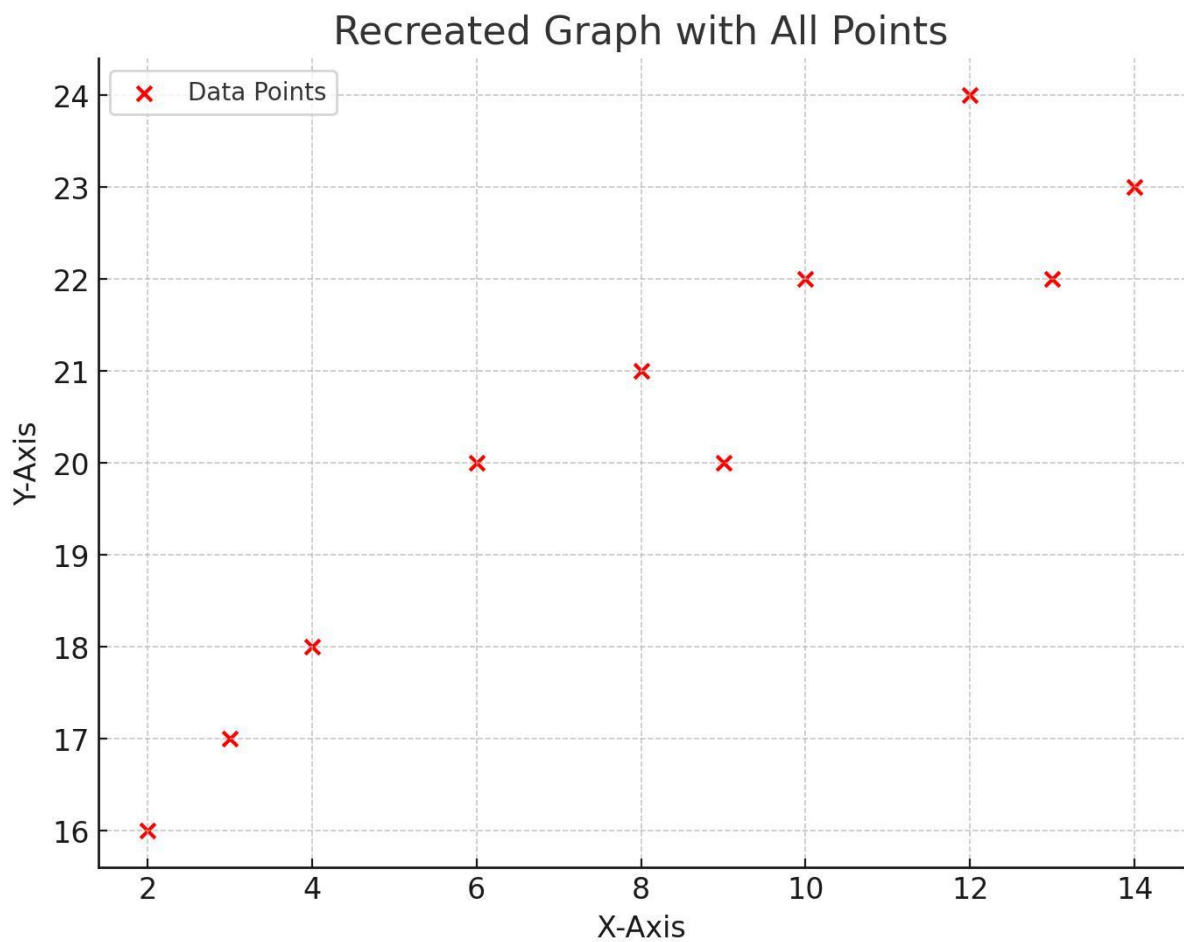
class

prediction = k.predict(np)

printf('Predicting the new point class')

print(prediction)

```
print('Plotting full data')  
plt.scatter(x = x + [new_x], y = y + [new_y], c = classes  
[prediction[0]], marker = '*')  
plt.show()  
OUTPUT:
```



EXPERIMENT: 3

AIM: Implement K mean clustering algorithm.

LIBRARIES IMPORTED: Numpy, Sklearn

INPUT:

```
import numpy as np
from sklearn.cluster import KMeans

#data points using numpy
data = np.array([
    [2, 1], [3, 4], [2, 4], [8, 8], [4, 1], [4, 3], [5, 6],
    [7, 8], [6, 5], [6, 2], [8, 1], [6, 1], [5, 2], [4, 2], [2, 5]
])

# Number of clusters (k)
k = 4

# Apply K-Means clustering
kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
kmeans.fit(data)
centroids = kmeans.cluster_centers_
labels = kmeans.labels_

print("Final Centroids:") #printing centroids
print(centroids)

print("\nClusters Formed:")
for i in range(k):
    # Grouping data points by clusters
    print(f"Cluster {i}: {data[labels == i]}")
```

OUTPUT SCREEN:

```
Final Centroids:
[[3.8         1.8         ]
 [6.5         6.75        ]
 [2.33333333  4.33333333]
 [6.66666667  1.33333333]]

Clusters Formed:
Cluster 0: [[2 1]
 [4 1]
 [4 3]
 [5 2]
 [4 2]]
Cluster 1: [[8 8]
 [5 6]
 [7 8]
 [6 5]]
Cluster 2: [[3 4]
 [2 4]
 [2 5]]
Cluster 3: [[6 2]
 [8 1]
 [6 1]]
>>> |
```


PRACTICAL 5


AIM: Linear Regression

LIBRARIES IMPORTED: Numpy, Sklearn

INPUT:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
np.random.seed(42)
x = 2 * np.random.rand(100, 1)
y = 4 + 3 * x + np.random.randn(100, 1)
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
random_state=42)
model = LinearRegression()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
print(f"Intercept: {model.intercept_}, Coefficient: {model.coef_}")
plt.scatter(x_test, y_test, color='blue', label='Actual Data')
plt.plot(x_test, y_pred, color='red', linewidth=2, label='Regression Line')
plt.xlabel("X")
plt.ylabel("Y")
plt.legend()
plt.show()
```

OUTPUT:



PRACTICAL 6

AIM: Logistic Regression

LIBRARIES IMPORTED: Numpy, Sklearn

INPUT:

```
import numpy as np

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

np.random.seed(42)

x = 3 * np.random.rand(200, 1) - 1.5 # Random values between -1.5 and
1.5

y = (x > 0).astype(int).ravel() # Label: 1 if x > 0, else 0

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
random_state=42)

model = LogisticRegression()

model.fit(x_train, y_train)y_pred = model.predict(x_test)accuracy =
accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)

report = classification_report(y_test, y_pred)


print(f"Accuracy: {accuracy:.2f}")

print("Confusion Matrix:")

print(conf_matrix)

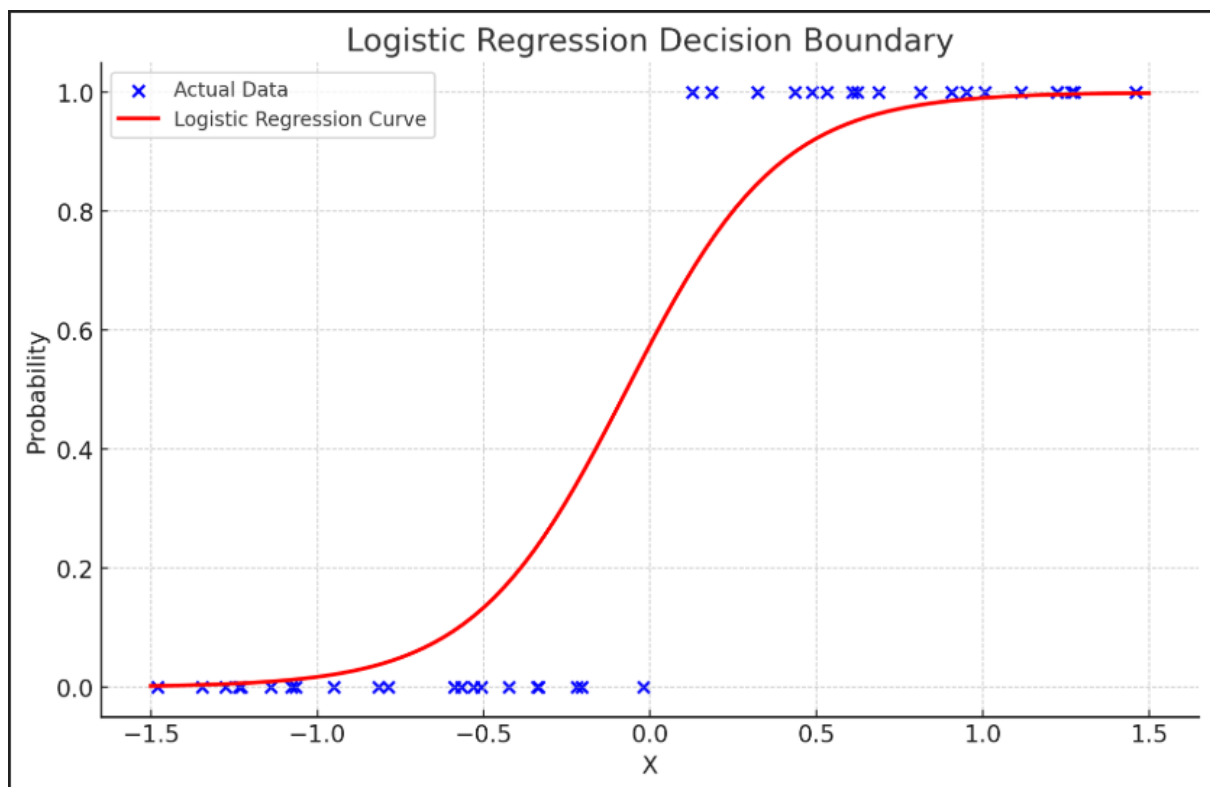
print("Classification Report:")

print(report)

x_values = np.linspace(-1.5, 1.5, 100).reshape(-1, 1)
```

```
y_prob = model.predict_proba(x_values)[:, 1]
plt.scatter(x_test, y_test, color='blue', label='Actual Data')
plt.plot(x_values, y_prob, color='red', linewidth=2, label='Logistic
Regression Curve')
plt.xlabel("X")
plt.ylabel("Probability")
plt.legend()
plt.title("Logistic Regression Decision Boundary")
plt.show()
```

OUTPUT:



	precision	recall	f1-score	support
0	1.00	0.95	0.98	21
1	0.95	1.00	0.97	19
accuracy			0.97	40
macro avg	0.97	0.98	0.97	40

PRACTICAL 7

AIM: Support Vector Machine

LIBRARIES IMPORT: Numpy, Sklearn

INPUT:

```
import numpy as np

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

np.random.seed(42)

x = np.random.randn(200, 2) # 2D data for visualization

y = (x[:, 0] * x[:, 1] > 0).astype(int) # Label: 1 if x1*x2 > 0, else 0
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
random_state=42)
model = SVC(kernel='linear')

model.fit(x_train, y_train)

y_pred = model.predict(x_test)

accuracy = accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)

report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")

print("Confusion Matrix:")

print(conf_matrix)

print("Classification Report:")

print(report)

def plot_decision_boundary(model, X, y):

    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
```

```


y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100), np.linspace(y_min,
y_max, 100))

Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
plt.contourf(xx, yy, Z, alpha=0.3)
plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k')
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.title("SVM Decision Boundary")
plt.show()

```

```
plot_decision_boundary(model, x_test, y_test)
```

OUTPUT:



	precision	recall	f1-score	support
0	0.57	0.89	0.69	19
1	0.80	0.38	0.52	21
accuracy			0.62	40
macro avg	0.68	0.64	0.61	40

PRACTICAL 8

AIM: Wine Quality Prediction

LIBRARIES: Pandas, Numpy

INPUT:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn import metrics
from sklearn.svm import SVC
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
import warnings
warnings.filterwarnings('ignore')
```

OUTPUT:

	type	fixed acidity	volatile acidity	citric acid	residual sugar	\
0	white	7.0	0.27	0.36	20.7	
1	white	6.3	0.30	0.34	1.6	
2	white	8.1	0.28	0.40	6.9	
3	white	7.2	0.23	0.32	8.5	
4	white	7.2	0.23	0.32	8.5	

	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	\
0	0.045	45.0	170.0	1.0010	3.00	
1	0.049	14.0	132.0	0.9940	3.30	
2	0.050	30.0	97.0	0.9951	3.26	
3	0.058	47.0	186.0	0.9956	3.19	
4	0.058	47.0	186.0	0.9956	3.19	

	sulphates	alcohol	quality
0	0.45	8.8	6
1	0.49	9.5	6
2	0.44	10.1	6
3	0.40	9.9	6
4	0.40	9.9	6